

Here is a technical summary of the challenges we faced during the development of **Chronosearch v2/v3**. These are excellent talking points for an interview because they show you can debug complex "Full Stack + Cloud" issues.

## 1. The "Zombie Code" Deployment

- **The Symptom:** You made changes to the code (like fixing the password length), but the error (**72 bytes limit**) kept happening exactly as before.
- **The Root Cause:** We were using a script (`setup_cloud.py`) that defined the Modal app but didn't explicitly "mount" the local `backend` folder. Modal was running a **cached, stale version** of the code from weeks ago, ignoring your local updates.
- **The Fix:** We switched to **Module-Based Deployment**.
- Command: `modal deploy -m backend.main`
- This forces Python to treat the folder as a package and syncs the exact local files to the cloud container.

## 2. The "72-Byte" Bcrypt Crash

- **The Symptom:** The backend crashed with `password cannot be longer than 72 bytes`, even when you manually sent a short password like "12345678".
- **The Root Cause:** It wasn't actually the password length. The library we were using, `passlib`, is old and unmaintained. It tried to read the version of the modern `bcrypt` library to check for compatibility but failed (`AttributeError: has no attribute '__about__'`). This crash triggered a misleading error message about password length.
- **The Fix:** We removed `passlib` entirely and rewrote `auth.py` to use \*\*pure `bcrypt`\*\*. This eliminated the dependency conflict and the error.

## 3. The "404 Not Found" / Connection Failed

- **The Symptom:** The frontend said "Connection Failed," and the Modal logs showed **zero** activity (no function calls).
- **The Root Cause:** A classic **API Mismatch**.
- **Frontend (api.js)** was sending a POST request to `/login` with **JSON** data.
- **Backend (api.py)** was listening on `/token` and expecting **Form Data** (`OAuth2PasswordRequestForm`).
- Because the URL and Content-Type didn't match, FastAPI rejected the request before it even entered your function code.
- **The Fix:** We renamed the backend endpoint to `/login` and updated it to accept a Pydantic model (`UserAuth`) so it could read the JSON.

## 4. The **ModuleNotFoundError** in Cloud

- **The Symptom:** The app deployed successfully, but as soon as a request hit it, it crashed with `No module named 'backend'`.

- **The Root Cause:** Python's "Relative Imports" (e.g., `from .common import app`) depend on where the script is executed from. In the cloud, the file structure didn't match the local environment's "root" context.
- **The Fix:** We standardized the imports to be absolute (`from backend.common import app`) and ensured the `modal deploy` command was run from the root directory so the package structure was preserved in the cloud.

## 5. Google OAuth "Origin Not Allowed"

- **The Symptom:** The Google Login popup appeared but immediately showed an error: `The given origin is not allowed for the given client ID.`
- **The Root Cause:** Google's security policy requires you to explicitly whitelist every domain that can initiate a login. `localhost:5173` was not on the list.
- **The Fix:** We updated the **Google Cloud Console** credentials to include `http://localhost:5173` in the "Authorized JavaScript Origins."

## 6. The "Infinite Indexing" Loop (Design Flaw)

- **The Symptom:** Every time you opened a video to watch it, the "Indexing" process would start again.
- **The Root Cause:** The `VideoPlayer` component (or the `upload` flow) didn't have a clear separation between "**Process Video**" (Write) and "**Watch Video**" (Read). The player was triggering the heavy AI job instead of just checking if the vectors already existed.
- **The Fix (Planned):** We are implementing a "**Write Once, Read Forever**" architecture where the Upload triggers the AI, and the Player only queries the database status.